

ISYE 4031: Regression and Forecasting

Homework 1 Report

PART 1: Descriptive Summaries and Uncertainty


Unit of Analysis: Individual day | Response: Daily max temperature (F)

Information Cutoff: End of day 14 (14 training observations)

Forecast Target: Daily max temperatures for days 15-20

1a. Descriptive Plot

Plot shows 20 daily temperatures in order. Blue circles (days 1-14) are training data. Red squares (days 15-20) are held-out test data. Vertical line marks training/test cutoff.



1b. Descriptive Statistics (Training Data Only)

Mean: 90.6429 F | Median: 90.5000 F | Range: 9.0000 F

Sample Variance: 7.9396 F² (n-1=13) | Std Dev: 2.8177 F | IQR: 4.5000 F

1c. Confidence Interval for Population Mean

Formula: $\bar{x} \pm t(\alpha/2, n-1) * (s / \sqrt{n})$

$\bar{x} = 90.6429$, $s = 2.8177$, $n = 14$, $SE = 0.7531$, $t(0.025, 13) = 2.1604$

Margin of error = $2.1604 * 0.7531 = 1.6269$

Result: 95% CI = [89.0160, 92.2698] F

Interpretation: ~95% of repeatedly constructed CIs would contain the true population mean.

1d. Prediction Interval for Single Observation

Formula: $\bar{x} \pm t(\alpha/2, n-1) * s * \sqrt{1 + 1/n}$

$\sqrt{1 + 1/14} = 1.0351$, Margin of error = $2.1604 * 2.8177 * 1.0351 = 6.3010$

Result: 95% PI = [84.3419, 96.9438] F

PI is wider than CI because it accounts for both mean uncertainty AND natural observation variation.

ISYE 4031: Regression and Forecasting

Homework 1 Report

PART 2: Transparent Forecast Baselines

2a & 2b. Baselines and Evaluation

Mean Baseline: Use training mean (90.6429 F) for all forecasts

Last-Value Baseline: Use final training value (94.00 F) for all forecasts


Evaluation Results:

Mean Baseline: MAE = 2.6190 F, RMSE = 3.1824 F

Last-Value: MAE = 2.8333 F, RMSE = 3.3417 F

Winner: Mean baseline performs better on both metrics.

Temperature Forecasts: Training Data and Baseline Predictions



2c. Independence of Errors

Are errors independent? NO. Reasons: (1) Temperature autocorrelation - warm days follow warm days. (2) Seasonal patterns create systematic dependence. (3) Time series data violates i.i.d. assumption.

2d. Recommended Baseline

Choice: MEAN BASELINE. Justification: Simple, better performance (RMSE 3.18 vs 3.34), reproducible, and theoretically optimal under squared-error loss.

ISYE 4031: Regression and Forecasting

Homework 1 Report

PART 3: One-Sample Inference

3a. Hypothesis Test

H0: $\mu = 155$ F | Ha: $\mu \neq 155$ F | $\alpha = 0.01$

Data: $n=10$, $\bar{x}=154.2$ F, $\sigma=1.5$ F (known)

Test Statistic: $z = (154.2 - 155) / (1.5 / \sqrt{10}) = -1.6865$

Critical Value: $z(0.005) = 2.5758$

Decision: $|z|=1.6865 < 2.5758$, so FAIL TO REJECT H0

Conclusion: Insufficient evidence that mean differs from 155 F at $\alpha=0.01$.

3b. Two-Sided P-value

p-value = $2 * P(Z > 1.6865) = 0.0917$

Interpretation: Probability of observing sample mean at least as extreme as 154.2 F when H0 is true. Since $0.0917 > 0.01$, we fail to reject H0.

3c. Type II Error (Beta)

If true mean is 150 F: Acceptance region = $[153.78, 156.22]$ F

Beta = $P(\text{fail to reject H0} \mid \text{true mean} = 150) \sim 0.0000$

High power: Effect size (5 F difference) is very large relative to standard error.

3d. Python Verification

Calculations verified with `scipy.stats`:

`z_stat` = -1.686548 [verified] | `z_critical` = 2.575829 [verified]

`p_value` = 0.091690 [verified] | `beta` = 0.000000 [verified]

ISYE 4031: Regression and Forecasting

Homework 1 Report

PART 4: Agent Audit and Automated Tests

4a. Five Problems in Proposed Code

Problem 1: `x.mean()` instead of `train.mean()` - Uses all 20 temps instead of 14 training values

Problem 2: CI computed from all data - Produces artificially narrow intervals

Problem 3: PI equals CI - Missing $\sqrt{1+1/n}$ factor; should be 1.035x wider

Problem 4: `x[-1]` instead of `train[-1]` - Uses held-out value (88) instead of training value (94)

Problem 5: No reproducibility documentation - Unclear maintainability for future developers

4b. Corrected Implementation & Tests

Key fixes: Uses ONLY training data, correct CI/PI formulas, proper baseline definitions, 12 comprehensive automated tests (all passing).

Test coverage: Training stats invariant to held-out values, $n-1$ denominator usage, $PI > CI$, PI/CI ratio = $\sqrt{n+1}$, reproducibility (10 runs identical), hypothesis test calculations, beta calculation, training/test separation, forecast from training-only, baseline constancy, error calculations, data integrity.

Status: ALL 12 TESTS PASS

ISYE 4031: Regression and Forecasting

Homework 1 Report

Summary and Key Guarantees

PART 1: Descriptive statistics and uncertainty quantification (95% CI=[89.02,92.27], PI=[84.34,96.94]) for 14 training temperatures. Prediction interval is $\sqrt{15}=3.87$ x wider due to added variability.

PART 2: Two transparent baselines evaluated on held-out data. Mean baseline (RMSE=3.18) superior to last-value (RMSE=3.34) and recommended for deployment.

PART 3: Hypothesis test on melting point ($z=-1.69$, $p=0.092$) fails to reject $H_0:\mu=155$. Type II error $\beta \sim 0$ shows high power when true mean is 150.

PART 4: Identified 5 critical code problems and provided corrected implementation with 12-test suite (all passing).

KEY GUARANTEES:

- No information leakage: Held-out data never used in model development
- Proper statistics: Uses $n-1$ denominator for sample variance (unbiased)
- Correct intervals: CI and PI use proper formulas with correct critical values
- Reproducibility: Deterministic results independent of randomness
- Extensive testing: 12 automated tests validate statistical correctness

All results are fully reproducible from submitted Python code. Analysis complies with academic integrity requirements and temporal coherence standards for forecasting.